Question 1 (basic): Complexity

The following function computes the *composite* of two maps.

```
function compose(p: Map<X, Y>, q: Map<Y, Z>) → Map<X,  Z>:

    r = new map with keys X and values Z

    for each key x in p:

        y = p.get(x)

        if q.containsKey(y):

            z = q.get(y)

            r.put(x, z)

    return r
```

All maps here are implemented using AVL trees. Assume that comparing two keys takes $O(1)$ time.

Let n be a number such that p and q have at most n elements.
What is the asymptotic (time) complexity of compose(p, q) in n?

Write your answer in O-notation. Be as exact and simple as possible.

Explain why the complexity of the function has this order of growth.
You can also add notes directly in the code above.

Answer

Complexity: $O(n \log(n))$

Justification:

Since p has at most n elements, the loop has n iterations. In each iteration, we put at most one element into r. Therefore, r also has at most n elements throughout the program. Since all of p, q, r have at most n elements, the calls p.get, q.containsKey, q.get, and r.put all take $O(\log(n))$. So the body of the loop takes $O(\log(n))$, meaning the whole function takes $O(n \log(n))$ time.

Question 2 (basic): Sorting

We run quicksort to sort the following sequence of letters alphabetically:

E	C	B	F	G	A	D

For partitioning a sequence S of values, we use the following simplified algorithm:

- The pivot is the **first element** (and does not belong to any part of the partition).
- For each part, we take the relevant elements from S in the **same order** as they occur in S.

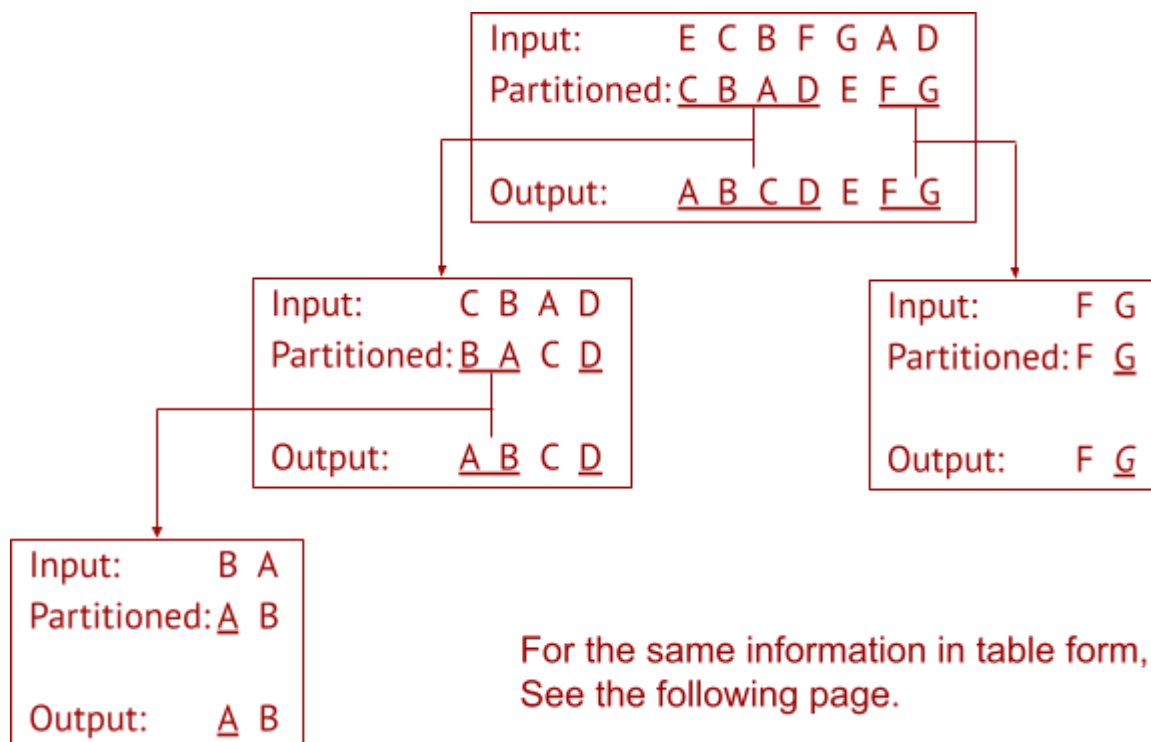For each call of quicksort, show:

- the sequence of values it is given (only the subarray it operates on),
- how the sequence looks after partitioning,
- what recursive calls are made,
- the sequence it returns.

Only show this when the sequence has **at least two elements**.
That is, ignore calls of quicksort on zero or one values.

Answer

Choose how you want to organise the requested information. For example, it could be a call graph (with arrows for recursive calls). You can also just use the answer table on the next page.


```
Input:      E  C  B  F  G  A  D
Partitioned: C  B  A  D  E  F  G

Output:      A  B  C  D  E  F  G
```

```
Input:      C  B  A  D
Partitioned: B  A  C  D

Output:      A  B  C  D
```

```
Input:      F  G
Partitioned: F  G

Output:      F  G
```

```
Input:      B  A
Partitioned: A  B

Output:      A  B
```

For the same information in table form,
See the following page.

Question 2: Optional answer sheet

For the recursive calls, you can underline each subsequence where a recursive call is made.

Input sequence	After partitioning	Recursive calls	Output sequence
E C B F G A D	C B A D E F G	C B A D E F G	A B C D E F G
C B A D	B A C D	B A C D	A B C D
B A	A B	A B	A B
F G	F G	F G	F G

Question 3 (basic): Lists, stacks, queues

The following function modifies a list of even size.

function twister(l: list):
      s = new stack (using a linked list)
      q = new queue (using a circular dynamic array)

      while l not empty:
          s.push(l.removeFirst())   push(A), push(B), ..., push(G), push(H)

      while s not empty:
          q.enqueue(s.pop())   enqueue(H), enqueue(G), ..., enqueue(B), enqueue(A)

      while q not empty:
          x = q.dequeue()
          y = q.dequeue()
          s.push(y)
          s.push(x)   push(G) & push(H), ..., push(A) & push(B)

      while s not empty:
          l.addLast(s.pop())   addLast(B), addLast(A), ..., addLast(H), addLast(G)

We call the function on the following list:

A	B	C	D	E	F	G	H

Answer

How does the list look afterwards?

B	A	D	C	F	E	H	G

Each of the two passes through the stack inverts the sequence. Everything else preserves the sequence, except for the dequeuing loop. There, we swap every group of two array values in the inverted sequence, but that is the same as swapping each such group in the original sequence.

Question 4 (basic): Hash tables

The below open-addressing hash table uses linear probing and modular compression (mod 10).

0	1	2	3	4	5	6	7	8	9
A	B	C			D	E			F

Unfortunately, I forgot the hash codes of all the elements.
I only remember they were non-negative and smaller than 20.

For each element, determine *all* the possible hash codes.

Answer

A) 0, 9, 10, 19

B) 0, 1, 9, 10, 11, 19

C) 0, 1, 2, 9, 10, 11, 12, 19

D) 5, 15

E) 5, 6, 15, 16

F) 9, 19

For a given element, the possible table indices are the slots in the same cluster until the element. This gives all the possible hash codes smaller than 10. For the possible hash codes below 20, we just add 10 to all the possible table indices (because of modular compression.)

Question 5 (basic): Search trees

The drawing on the right *potentially* depicts a red-black tree (null nodes omitted).

First, decide if it is indeed a red-black tree.

- If your answer is **negative**, explain what is wrong.
  State the relevant part of the invariant and where in the tree it is broken.

- If your answer is **positive**, insert 5 into it.
  State the rebalancing steps and show the resulting red-black tree with red nodes marked.

  If you prefer, you can use the insertion procedure for 2-3 trees. In that case,
  also show the conversion to a 2-3 tree and the conversion back to a red-black tree.

Note: if you use a version of red-black trees different from your course, you must define it.

Answer

Is this a red-black tree?   ☑ Yes   ☐ No

Second task:

Binary search puts the new red node for 5 as the right child of 2.
Rebalancing steps:

1. Left rotate the edge 2–5.
2. Right rotate the edge 8–5.
3. Perform a colour flip of 5 and its children.
4. Colour the root black.

Here is a solution using 2-3 trees:

1. The starting tree corresponds to a 2-3 tree with just a single 3-node: 2|8.
2. Inserting 5 and merging upwards creates a 4-node, which we split as two levels of 2-nodes: 5 with left child 2 and right child 8.
3. We are left with only 2-nodes, so the corresponding red-black tree is identical, with all nodes coloured black.

Question 6 (basic): Priority queues

You are given the following binary heap.

| 0	1	2	3	4	5	6	7	8
19	41	23	68	99	41	30	75	72

Answer

Is this a min-heap or a max-heap? min-heap

Draw the heap in tree representation:

Add 36 to the heap. What is the tree representation now?

Question 7 (basic): Graphs

Perform a uniform-cost search (Dijkstra's algorithm) in the weighted graph to the right, starting from **G**.

Trace your steps in the table below. Write priority queue entries as X:c where X is the graph node and c is its cost.

Time-saving instructions:

- Do not add an entry to the priority queue if its node is already visited.
- When comparing two entries with the same cost, compare their nodes **alphabetically**.
- We are only interested in calculating distances here, not actual full shortest paths.

Answer

PQ entry to process	Visit? yes/no	Priority queue (after processing entry)
—		G:0
G:0	yes	A:1 E:2
A:1	yes	B:2 E:2 C:3
B:2	yes	E:2 C:3 H:3 D:4
E:2	yes	C:3 H:3 D:4 C:6
C:3	yes	H:3 D:4 C:6 D:7
H:3	yes	D:4 F:5 C:6 D:7
D:4	yes	F:5 C:6 D:7 F:8
F:5	yes	C:6 D:7 F:8
C:6	no	D:7 F:8
D:7	no	F:8
F:8	no	

Question 8 (basic): Abstract data types

While surfing the web for interesting new data structures to learn, you suddenly stumble upon the following obscure code:

class Thermometer
      hot: Thermometer
      cool: Thermometer
      fridge: integer
      oven: string

function temperature(t: Thermometer, x: integer) → string:
      if t is null:
            return null
      if x < t.fridge:
            return temperature(t.hot, x)
      if x > t.fridge:
            return temperature(t.cool, x)
      return t.oven

At first, this looks very intriguing, but then you realise that you already know this data structure!

Answer

Which data structure is this? A binary search tree

Find better names (in the context of that data structure) for the following:

      class Thermometer = Node or BSTNode or BST

      function temperature = get or lookup

Question 9 (advanced): Complexity

There are *N* cities with populations. You are given the following sorted arrays of length *N*:

- The array **cities** stores only the city names, but is sorted by *population*.
- The array **populations** stores the pairs of a city name and its population, and is sorted alphabetically by *name* (the first component of the pair).

Here is an example:

	cities		**populations**
0	Tallinn	0	(Amsterdam, 921402)
1	Bratislava	1	(Athens, 637798)
2	Lisbon	2	(Berlin, 3677472)
3	Vilnius	3	(Bratislava, 475044)
4	Dublin	4	(Budapest, 1706851)
	. . .		. . .
N−2	Madrid	*N*−2	(Stockholm, 978770)
N−1	Berlin	*N*−1	(Tallinn, 437811)

Consider the following problem:

"Given a number *K*, what is the name of the largest city with population less than *K*?"

We want to solve this problem using a fast algorithm. What is the best asymptotic complexity in *N* you can achieve? Write your answer in *O*-notation. Justify by briefly sketching your algorithm.

- Treat number and string comparisons as taking constant time.
- You do not have to justify why there is no better algorithm.

Answer

Complexity: $O(\log(N)^2)$

Justification: Note first that the array **populations** can serve as an immutable map with city names as keys and populations as values. We use binary search to look up the population for a city name ($O(\log(N))$).

To find the largest city with population less than *K*, we perform a binary search in the array **cities**. Each step uses a comparison of the current city's population with *K*, which we said is $O(\log(N))$. Since there are $O(\log(N))$ many steps in binary search, the total runtime is $O(\log(N)^2)$.

(Pedantic remark: since two cities may have the same population, we may want to find *all* largest cities with population less than K, not "a largest city". For this, we can run the version of binary search that finds the first and the last matching element. This also takes $O(\log(N)^2)$.)

Question 10 (advanced): Hash functions

Let *T* be a type (e.g., a class) with *N* different possible values (e.g., objects). We call a hash function for *T perfect* if it has **no collisions** and its values are the **first *N* natural numbers** (0, …, *N*-1).

Perfect hash functions are useful because they allow a very simple hash table design with good worst-case behaviour.

Here is my specification of a student:

> **class** Student:
>   grade: one of U, 3, 4, 5
>   height: integer from 110 to 210 (that is, 110, 111, …, 209)
>   nice: boolean (true or false) indicating if the student is nice

Design a perfect hash function for this type. It must run in constant (O(1)) time.

Notes:

- Students with the same grade, height, and niceness are considered equal.
- You may answer using pseudocode or your favourite programming language.

Answer

Number of different possible students (*N*): $4 \cdot 100 \cdot 2 = 800$

There are many ways to build a perfect hash function in this example.

A modular strategy is to build a perfect hash function for each attribute and then combine them using basic combinatorics. In this case, there are obvious perfect hash functions for the attributes:

- hash_grade: map U, 3, 4, 5 to 0, 1, 2, 3 (in this order, *N* = 4)
- hash_height: subtract 110 (*N* = 100)
- hash_nice: map true to 1 and false to 0 (*N* = 2)

One systematic way of combining these perfect hash functions is as follows:

function perfect_hash(s: Student) → integer:
    h = 0
    h = 4 · h + hash_grade(s.grade)
    h = 100 · h + hash_height(s.height)
    h = 2 · h + hash_nice(h)
    return h

This enumerates all the possible students in *lexicographic order* of (grade, height, nice). That is, first all the students with grade U, then all those with grade 3, then grade 4, then grade 5. Within each grade, first all students with height 110 (first the nice one, then the nice one), then height 111, ….

Question 11 (advanced): Linear search trees

A tree is called *linear* if every node has at most one child that is not null.
Equivalently, its height is equal to its size.

We insert the natural numbers below N (that is, 0, ..., $N-1$) into
an empty binary search tree (BST). Depending on the order
of insertions, we may end up with a linear tree.

For example, for $N = 5$, inserting in the order 4, 0, 1, 3, 2 creates
the linear BST depicted on the right (null nodes omitted).

As you know, the sequence of insertions takes $\Theta(N^2)$ time in that case.
We wish to detect that case (final BST is linear) more efficiently.

Write an algorithm that takes:

- a natural number N,
- an array x of length N containing the natural numbers below N in some order

and returns a boolean. The result should be *true* exactly if inserting the values in x in that order into
an empty BST would create a linear tree.

Your algorithm must have (time) complexity $O(N)$.

You may answer using pseudocode or your favourite programming language.

Answer

One possible algorithm:

function test_linear(N: integer, x: array of integers) → boolean:
    lo = 0
    hi = N - 1
    for each k **in** x:
        if k = lo:
            lo = lo + 1
        else if k = hi:
            hi = hi - 1
        else:
            return false
    return true

The problem becomes a bit harder if the set of elements in x is not known in advance.
That is, only the length of x is given. Can you figure out how to solve the problem then?

Question 12 (advanced): The key master

You have been taken prisoner by the key master! To escape, you need to collect all N keys, conveniently numbered in sequence (0, 1, ..., N–1).

The key master has sealed each key behind a gate. That gate **may** need another key to open. After opening it, you can collect the key behind it.

Clarifications:

- The keys you collect stay with you forever. They do not get used up by opening gates.
- The same key may be needed to open several gates.

Design an algorithm to help you escape. The input consists of:

- the number N of keys,
- an array x of length N where $x[k]$ is the key needed to collect key k or *null* if no key is needed.

Print the keys in an order in which you can collect them, or raise an error if this is impossible.

But beware! If your algorithm does not run in $O(N \log(N))$ time, you will starve to death.

Notes:

- You may use any data structure or algorithm from the course in your solution. You don't need to implement those yourself.
- You may answer using pseudocode or your favourite programming language. Pseudocode can also include clear and precise natural language description.

Answer

This is an instance of a *scheduling problem* (task $x[k]$ must come before task k). There are many possible solutions here, and most are actually $O(N)$. The additional factor of $\log(N)$ is just to allow uses of search-tree based maps or sorting-based approaches if desired.

Many solutions are based on the following two ingredient:

- An auxiliary "index" map sending each key to the list of keys it unlocks. Since our keys are indexed sequentially, we can implement this index using an array. Building it takes $O(N)$.

- A queue of the keys that we can collect using only the keys we already have. It starts with all the keys that do not need a key to be collected. We loop over this queue as our agenda. Every time we process a queue entry, we collect the associated key and add to the queue all the keys that this key unlocks. This loop is $O(N)$.

In this approach, we need to check after the queue is empty if we managed to collect all the keys. If not, we are stuck forever and raise an error.

Example code is given on the next page.

```
function escape(N: integer, x: array of "integer or  null"):
        index = new array of lists of integers (initially all empty)
        can_collect = new queue of integers
        for k from 0 to N:
                if x[k] is null:
                        can_collect.add(k)
                else:
                        index[x[k]].add(k)

        while can_collect is not empty:
                k = can_collect.remove()
                "collect key k"
                for m in index.get(k):
                        x[m] = null
                        can_collect.add(m)

        if x contains non-null: raise error "trapped forever!"
```

It is possible to avoid an index by inverting the control flow: start with a locked key and recursively collect any prerequisite key as a pre-processing step. In this approach, we need to raise an error if a dependency cycle is encountered. To detect this, we maintain a set of "ancestors" to remember all calls in the recursion we are currently processing.

```
function escape(N: integer, x: array of "integer or  null"):
        ancestors, collected = new sets of integers

        function collect(k: integer):
                if k not in collected:
                        if k in ancestors: raise error "trapped forever!"
                        with k added to ancestors (removing it afterwards):
                                If x[k] is not null: collect(x[k])
                                "collect key k"
                                collected.add(k)

        for k from 0 to N: collect(k)
```

In fact, you may recognize this algorithm: it is depth-first search (DFS) in the graph that has keys as nodes and an edge $m \to k$ if m locks the gate for k. Specifically, it is the version of DFS that produces a (reversed) topological sort!

So a high-level, one paragraph answer is: perform a topological sort of the key dependency graph encoded by x. A topological sort, if it exists, is the reverse order in which to collect the keys.

By the way: all of the above solutions also work if each key is locked by a *sequence* of gates. Then the runtime is linear in the number of keys plus the number of gates.