

Basic question 1: Sorting

Here is a sketch of quicksort:

```
function qs(array: array of integers, low: integer, high: integer):  
    if low < high:  
        pivotPosition = partition(array, low, high)  
        qs(array, low, pivotPosition-1)  
        qs(array, pivotPosition+1, high)
```

The function call `partition(array, low, high)` chooses a pivot in an unknown way and partitions the array range `[low, high]` before returning the new position of the pivot element.

Consider the following array:

0	1	2	3	4	5	6	7	8
10	5	7	3	9	2	8	3	11

Calling `partition(array, 0, 8)` returns 4.

- What was the chosen pivot value?
- Which values do the parts of the partition now contain?

Note: the order of values in your answers does not matter.

Pivot value: _____

Values in left part: _____

Values in right part: _____

Basic question 2: Hash tables

The following open addressing hash table models a set S of animals. It uses modular compression (using the modulo operator) and linear probing (with probing constant 1, as usual).

0	1	2	3	4	5	6	7	8	9
Fly	Cat			Bee				Gnu	Ant

For each of the following animals x , state how many array cells (possibly empty) the call

$S.contains(x)$

accesses.

Animal	Hash code	Number of array accesses
Bee	7114	
Cat	3650	
Dog	9374	
Elk	1509	

Basic question 3: Search trees

Draw an AVL tree representing the set of integers $\{1, 2, 3, 4, 5, 6, 7\}$ that is as **unbalanced** as possible (that is, has maximum height without breaking the AVL invariants).

Basic question 4: Priority queues

The following array represents a binary max-heap:

0	1	2	3	4	5	6	7	8	9
23	14	20	11	8	3	7	2	7	5

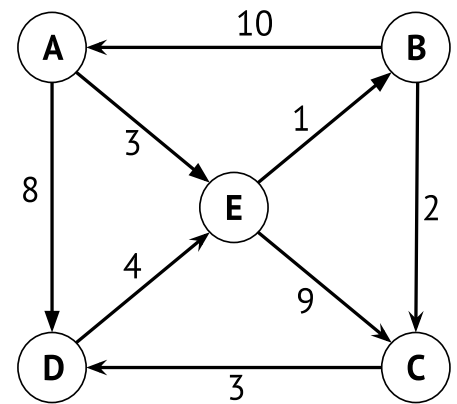
- Draw the tree representation of the heap.
- Draw the tree representation of the heap after adding the value 21 to it.

Basic question 5: Graphs

We run this familiar piece of code on the graph to the right:

```

visited = new map from nodes to distances
agenda = new min-priority queue of pairs of cost and node
agenda.add( (0, A) )
while agenda is not empty:
    (cost, node) = agenda.removeMin()
    if not visited.containsKey(node):
        visited.put(node, cost)
        for e in outgoingEdges(node):
            agenda.add( (cost + e.weight, e.target) )
    
```



We run three iterations of the while-loop and then stop. **At that point:**

- What are the three entries of the *visited* map (you get one for free)?

A \mapsto 0

_____ \mapsto _____

_____ \mapsto _____

- What are the four entries of the *agenda*, in ascending order?

Cost	Node

Basic question 6: Complexity

Your boss wants you to calculate the product $a \cdot b$ for large integers $a, b \geq 0$. Just after your multiplication operator broke down last week! Here are two workarounds you came up with:

```
function mulA(a, b):  
    if a == 0:  
        return 0  
    return b + mulA(a - 1, b)
```

```
function mulB(a, b):  
    result = 0  
    while a > 0:  
        if a % 2 == 1:  
            result = result + b  
        result = result + result  
        a = a // 2  
    return result
```

State the asymptotic time complexity of each algorithm **in terms of a** .

In each case, **briefly state** how you concluded this (you may do this by annotating the programs).

Notes:

- Assume that all arithmetic operations take $O(1)$ time.
- If you answer in O -notation, be as exact and simple as possible.
- We write $a // d$ for the *quotient* and $a \% d$ for the *remainder* in integer division of a by d .

Asymptotic complexity of mulA: _____

Short justification:

Asymptotic complexity of mulB: _____

Short justification:

Advanced question 7: Heapification

Recall the helper functions that we use to implement a binary heap:

- *swim* (also called *swimming up* or *sifting up*),
- *sink* (also called *sinking down* or *sifting down*).

Here are two algorithms for turning an array of integers into a binary heap:

- (1) Go over the array in *forward* order and call *swim* at every cell.
- (2) Go over the array in *backward* order and call *sink* at every cell.

For each algorithm, determine and justify the asymptotic time complexity in the array size n . Based on your analysis, conclude which algorithm has better scaling behaviour.

Hint: In a complete binary tree, the average node height is $O(1)$. You can use this without proof.

Advanced question 8: Water world

In the game of water world, a player must navigate a hexagonal grid of map tiles without drowning. The grid is implemented using this class:

```
class Tile:                                // can be compared/hashed in O(1) time
    elevation: number
    neighbours: list of Tile  // at most six neighbours
```

The player can jump from a tile to a neighboring tile if their elevations differ by at most 2.

The water level of the world starts at 0 and increases by 1 after making a jump.
The player drowns if the water level exceeds the elevation of their tile.

Specify an algorithm

```
function solvable(start: Tile, goal: Tile) → boolean
```

that returns true if the player can get from the start tile to the goal tile without drowning.
Your algorithm should run in average $O(N)$ time where N is the number of grid tiles.

Notes:

- You do not have to justify the complexity of your algorithm.
- You can use **unambiguous** natural language or pseudocode to describe your algorithm.
- You can freely use data structures and algorithms from the course – you do not have to explain how they work. In particular, hash sets of tiles have average-case $O(1)$ operations.



Advanced Question 9: Ordered Set

Recall that a **set** of values of type T has the following operations:

- $\text{contains}(x: T) \rightarrow \text{boolean}$: check if x is in the collection.
- $\text{add}(x: T)$: add x to the collection unless it is already an element.

An **ordered set** additionally remembers the order in which its elements were added.

For our purposes, it should support the following additional operations:

- $\text{pop}() \rightarrow T$: remove and return the **newest** element of the ordered set.
This method assumes that the ordered set is not empty.
- $\text{replace}(x: T, y: T)$: replace x with y without changing its position in the order of elements.
This method assumes that x is an element and y is not currently an element.

We assume that T has a good, constant-time hash function. Design a data structure implementing an ordered set such that all of the above methods run in **average $O(1)$ time**. Your answer should:

- define the data structure,
- state the implementations of *pop* and *replace*.

Note: you can freely use data structures and algorithms from the course – you do not have to explain how they work.

Hint: A standard hash table implements a set with the desired complexity, but not an ordered set. A hash table can still be a useful ingredient of your solution (together with additional structure).